

The Operational Risk Audit

The 2025 Autumn Governance Series: Audit Committee Briefing on AI-Specific Op Risks



TANYA MATANDA

OCT 01, 2025



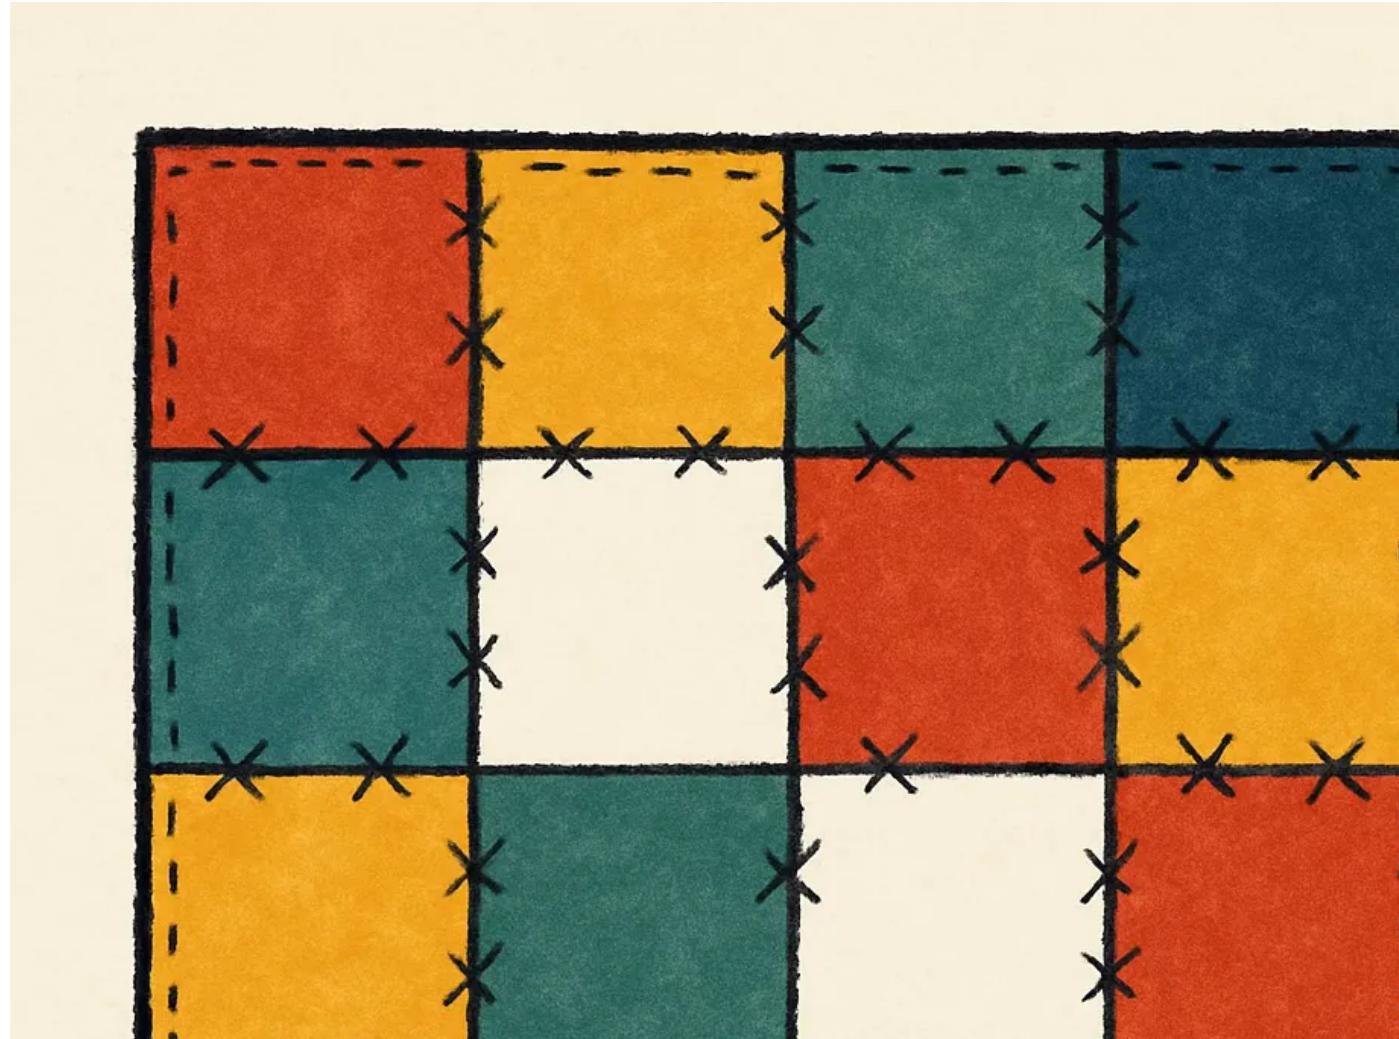
Sha

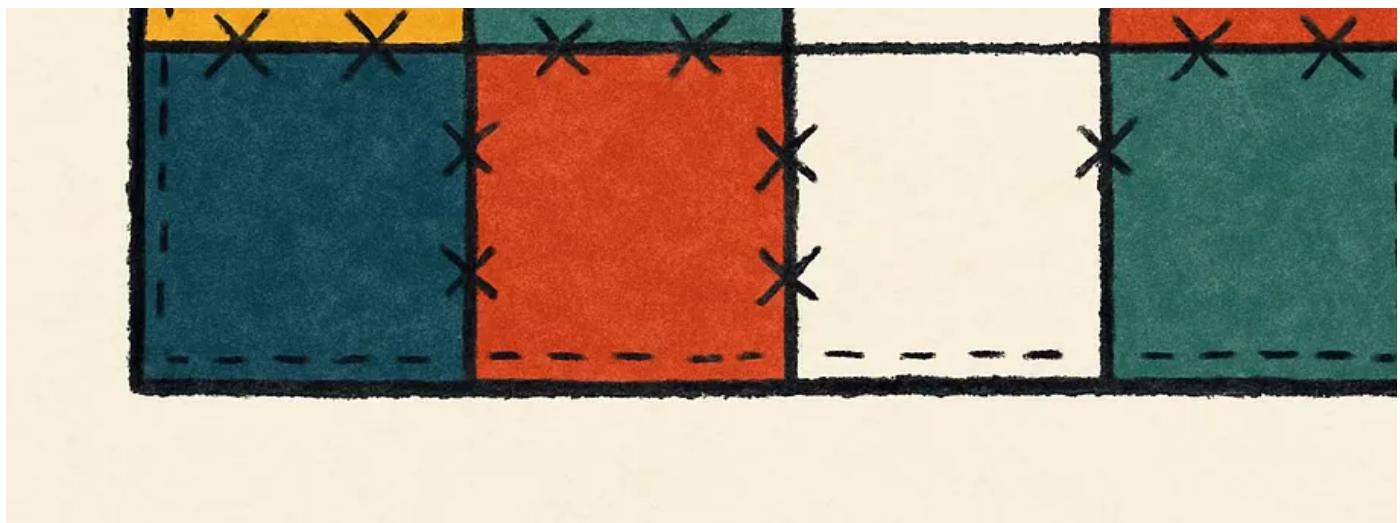


Article voiceover

0:00

-16:0





Artificial intelligence is transforming private equity, rapidly moving from experimental pilot projects to the backbone of investment analysis and portfolio company operations. Eighty-two percent of PE firms now rely on AI for core decisions such as deal sourcing, portfolio monitoring, and exit planning. Yet only one percent of those firms describe themselves as “operationally mature.” The gap between adoption and governance is wide – and boards are increasingly exposed.

This briefing analyzes the operational risks of AI, explains why traditional insurance policies provide limited protection, and offers concrete steps boards can take. Case examples illustrate how AI failures have already impacted valuations. A quantified assessment shows uninsured exposures could range from eight to thirty percent of fund value. The briefing closes with a roadmap for governance, investment and insurance strategy, balancing risk with opportunity.

Key Takeaway: AI governance should not be treated as a compliance burden. When done properly, it preserves value, reduces insurance costs, and strengthens fundraising.

Technology Failures and the Quality Crisis

The adoption of AI in private equity has outpaced its reliability. According to the Software Improvement Group (2024), three-quarters of AI systems suffer from severe quality issues, particularly in maintainability — the ability to update, repair, and safely extend systems. This matters because while equity funds often run on ten-year horizons, AI models can degrade within 12–18 months if not actively retrained.

This mismatch produces compounding risks. A degraded AI model may misprice secondary opportunities, overlook early warning signs in portfolio companies, or recommend investments in underperforming GPs. Each chip away at fund performance in ways that can be difficult to detect before the damage is already realized.

Case Example: In 2024, a mid-sized PE firm discovered during exit that its portfolio company's AI-driven sales recommendation engine had been degrading for years. The issue, identified by the buyer's diligence team, was expected to reduce valuation by 18 percent. The fund lost tens of millions, not from market conditions, but from an overlooked AI system.

Board Actions

- Require an **inventory of all AI systems** used at the fund and portfolio levels.
- Commission **annual maintainability scoring** of critical systems.
- Mandate reporting on **AI system degradation and retraining cycle**.

The Insurance Coverage Gap

Traditional insurance is designed for human error – random, diversifiable, actuarially predictable. AI liability does not fit these assumptions. Research (arXiv:2106.00839) highlights three differences:

1. **Concentration:** AI replaces multiple human decisions with one more concentrated risk.
2. **Unpredictability:** Models perform well on known data but fail suddenly on new patterns.
3. **Systematic Errors:** When AI fails, it tends to make the same mistake everywhere, amplifying loss.

Insurers are responding not by innovating coverage, but by excluding risk. Cyber policies exclude losses from model drift, adversarial manipulation, contaminated training data. D&O coverage is ambiguous on whether board members are liable for inadequate AI oversight. E&O increasingly excludes liability for AI-based recommendations.

Case Example: A U.S. venture fund reported to its LPs in 2024 that a cyber policy did not respond to losses after its AI monitoring system failed to detect fraudulent activity in a portfolio company. The insurer classified it as a “algorithmic error,” outside the scope.

Board Insurance Roadmap

1. **Inventory policies** (D&O, E&O, cyber, fiduciary).
2. **Audit exclusions** for algorithmic decision-making.
3. **Engage AI-specialized brokers and legal counsel.**

4. Quantify uninsured exposure using AI adoption data.
5. Explore alternatives: captives, parametric triggers, or government backstop programs (arXiv:2409.06672).

Source: Willis Towers Watson (2023–24), ABA reviews, arXiv:2106.00839.
Method: binary coding of coverage (0=excluded, 1=covered) across cyber, and E&O policies vs. five AI risk categories. As of Oct 2025.

Portfolio Company Operational Risks

One in five portfolio companies now embeds AI into operations, from forecasting to pricing. This increases efficiency but also creates dependencies that buyers scrutinize. Poor AI governance at the portfolio level translates directly into fund-level valuation risk.

Case Example: A European PortCo using AI for supply chain optimization achieved 15 percent cost reductions. But at exit, the buyer discounted valuation after discovering the system relied on a single vendor with no portability provisions. The result: a 22 percent haircut and an extended due diligence process.

Governance Standards for Portfolio Companies

- **Pre-Investment:** Assess AI system quality, vendor dependencies, and ownership.
- **During Hold:** Conduct quarterly AI performance monitoring and annual vendor risk reviews.
- **Pre-Exit:** Provide documentation of AI systems, vendor transferability, and regulatory compliance.

Cybersecurity as Financial Risk

AI introduces cyber risks that go beyond data breaches. Three stand out:

- **Model Theft:** Competitors steal proprietary models, erasing years of investment and eroding returns.
- **Training Data Contamination:** Attackers corrupt datasets, producing biased or erroneous results.

flawed outputs for months.

- **Algorithmic Manipulation:** Adversaries cause systems to produce decisions at scale.

Each risk carries exposures ranging from \$5 million to over \$200 million, limited to no insurance coverage. When AI controls physical systems (in logistics or manufacturing), cyber risks can morph into product recalls, incidents, or business interruption — typically excluded under both commercial property and liability insurance.

Board Actions

- Add **AI-specific metrics** to cyber dashboards (e.g., number of advanced attacks blocked).
- Require **annual penetration testing** of AI systems.
- Ensure **business continuity plans** include real-time decision contingencies, not just system recovery.

Vendor Concentration and Contingency Planning

AI vendor markets are highly concentrated. Three to five platforms dominate fund-level tools, while a handful of data providers supply benchmarks. A disruption—whether bankruptcy, security breach, or sudden contract termination—could ripple across the industry.

Case Example: In 2025, a U.S. analytics vendor discontinued a dataset and imposed secondary pricing. Several PE firms had to retrain models for \$2–3 million each, while performance degraded for nearly a year.

Board Actions

- Categorize vendors by criticality (Tier 1: irreplaceable, Tier 3: replaceable).
- Require contracts with **audit rights, data portability, and escrow clauses**.
- Test contingency plans with alternative vendors annually.

Business Continuity for Algorithmic Decision-Making

Traditional continuity planning focuses on restoring IT after outages. This requires a different lens: many investment decisions cannot wait. Deal opportunities, rebalancing, or crisis response often require action within hours.

Boards must therefore test continuity under degraded AI conditions. This includes simulating model drift, adversarial attacks, and vendor outages, not just server failures. Manual override protocols should be drilled, and cross-vendor redundancy established.

Board Actions

- Run **semi-annual drills** simulating AI failure during critical decision windows.
- Train staff in **manual override protocols**.
- Require **cross-vendor backups** for mission-critical processes.

Quantifying the Uninsured Exposure

Based on adoption rates and failure probabilities, uninsured AI exposure for a \$2 billion fund are estimated as:

- **Conservative:** \$165 million ($\approx 8\%$ of fund value).
- **Expected:** \$335 million ($\approx 16\%$ of fund value).
- **Severe:** \$630 million ($\approx 31\%$ of fund value).

For comparison, traditional operational insurance typically covers 1–3% fund value.

Source: V7 Labs PE AI Survey (Q4 2024), Software Improvement Group (Sergey et al., 2024), Bain & Co (2024), McKinsey (2025), arXiv:2106.00839 & 2409.06672. Methodology: This slide is based on a survey of 100+ companies across various industries. The survey asked about their use of AI for operational risk management. The results show that 80% of respondents believe AI has improved their operational risk management processes. The survey also found that AI is most commonly used for risk identification and monitoring, followed by risk assessment and mitigation. The survey also found that AI is most commonly used for risk identification and monitoring, followed by risk assessment and mitigation.

modelled loss estimates per \$2B fund across five risk categories; scenario lower/mid/upper bound assumptions. As of Oct 2025.

Governance Investment ROI

Governance is often seen as a cost center, but here it is demonstrably driver. A comprehensive AI governance program costs around \$1 million year, but delivers 5–15x ROI through:

- **Incident avoidance:** Preventing \$5–20 million losses per failure.
- **Performance protection:** Preserving 1–2% IRR.
- **Insurance optimization:** Cutting premiums by 15–30%.
- **Fundraising advantage:** Attracting LPs with stronger governance narratives.

Opportunity Framing: With robust governance, AI is not just safer – it becomes a competitive differentiator. Funds with credible governance frameworks can reassure LPs, negotiate better insurance, and realize returns from AI with fewer disruptions.

Source: EY AI Pulse Survey (2024), Private Funds CFO Forum (Feb 2026), McKinsey (2025). Method: scenario modelling of incident probability (3-y horizon) and uninsured exposure per \$2B fund under Minimal, Moderate, Comprehensive governance. As of Oct 2025.

Private Equity Audit & Risk Oversight GPT

A boardroom-ready governance co-pilot for private equity funds. It performs operational risk audits, insurance coverage analysis, portfolio oversight, risk quantification, vendor stress-tests, governance ROI, and fiduciary

checklists in one tool. [Link Here.](#)

Conclusion

AI has become integral to private equity, but the risks of poor governance are now unavoidable. Insurance gaps, vendor concentration, cybersecurity threats, and portfolio dependencies expose boards to hundreds of millions of uninsured risk. Yet the path forward is clear: invest in governance, embed operational risk oversight, and treat insurance as a dynamic strategy rather than a static policy.

The opportunity is significant. By moving early, boards can both reduce exposure and signal to LPs that they are leaders in responsible AI adoption. Governance becomes not just a shield, but a source of trust and value creation.

The insights from the Operational Risk Audit article resonate directly with governance frameworks I explore in *Shaping the Next Decade*. Both argue boards and committees must move beyond pedigree and static benchmarks learning to underwrite conviction edges while managing new risks introduced by AI, regulation, and systemic concentration. What begins as a framework for today's investment committees—scoring conviction and networks—becomes tomorrow's blueprint for fiduciary governance in an era where human judgment, algorithmic judgment, and long-horizon policy priorities intersect. Available for purchase here.

Research and Audio Supported by AI Systems



1 Like

← Previous

Discussion about this post

[Comments](#) [Restacks](#)



Write a comment...

© 2025 Tanya Matanda • [Privacy](#) • [Terms](#) • [Collection notice](#)
[Substack](#) is the home for great culture